

Yuanbo Pang

yuanbopang@berkeley.edu | 408-460-7526 | Berkeley, CA | [linkedin.com/in/yuanbopang](https://www.linkedin.com/in/yuanbopang) | github.com/PangYuanbo | yuanbopang.pro

UC Berkeley EECS student building AI agent evaluation, infrastructure, and robotics/vision systems; co-first author on the arXiv-published Agents' Last Exam benchmark.

SKILLS

AI/ML: agent evaluation, benchmark infrastructure, PyTorch, TensorFlow, CUDA, LLaVA, OpenCV | **Agent Systems:** agent harnesses, tool-calling workflows, VM/GUI/CLI evaluation, artifact capture, Claude Code, Codex, Cursor, Grok, Droid | **Engineering:** Python, TypeScript, FastAPI, React, Docker, Modal, REST APIs, GraphQL, WebSockets | **Vision/Robotics:** UAV detection, optical flow, ROS2, 3D perception

EDUCATION

University of California, Berkeley <i>B.S. Electrical Engineering and Computer Sciences (EECS)</i> Coursework: CS 61B, CS 70, CS 170, CS 184, EECS 16A, EECS 106A. Stanford Summer School 2025: CS 229 Machine Learning.	Berkeley, CA <i>Aug 2025 - Dec 2026</i>
De Anza Community College <i>Computer Science transfer coursework; Phi Theta Kappa; Academic Director, Quantum Club; President, Data Club</i>	Cupertino, CA <i>Sep 2023 - Jun 2025</i>

EXPERIENCE

Full-Stack Engineer, Ludus <i>AI agent marketplace, full-stack systems, agent execution</i> - Built outcome-first marketplace where multiple coding/research agents execute the same task in parallel and users pay only for accepted results. - Built React/TypeScript frontend, Node/Hono backend, Neon Postgres task store, and Docker-based VM dispatch layer for agent execution.	Remote <i>Apr 2026 - Present</i>
Undergraduate Researcher, UC Berkeley - Prof. Avideh Zakhor <i>UAV small-object detection, motion priors, airborne target tracking</i> - Research small UAV detection and tracking in aerial video, with focus on tiny-object recognition, camera-motion compensation, and low-false-positive perception for onboard drone systems. - Benchmarked UAV detection methods including TransVisDrone, AOT Winner v022, ESOD, and Li-TETC/NPS against 1,600 manually annotated DJI frames. - Evaluated optical-flow and background-compensation pipelines to reduce false positives in tiny airborne object tracking.	Berkeley, CA <i>Mar 2026 - Present</i>
ENGR Team Co-Lead, UC Berkeley BAIR - Agent's Last Exam (ALE) <i>AI agent evaluation, benchmark infrastructure, agent harnesses</i> - Co-first author on ALE, an arXiv-published benchmark led by Prof. Dawn Song for evaluating frontier AI agents on long-horizon professional workflows. - Primary builder for turning expert-submitted real-world tasks into executable benchmark instances across 55 sub-fields and 13 industry clusters. - Built agent evaluation infrastructure for VM-based GUI/CLI workflows, reproducible harness runs, and deliverable-based scoring.	Berkeley, CA <i>Jan 2026 - Present</i>
Robotics AI Engineering Intern, Starbot <i>Agent harness for restaurant service robots</i> - Built agent harness for restaurant service robots with tool APIs, task-state orchestration, and recovery paths for ordering, delivery, and customer-service workflows. - System was demonstrated at CES 2026 and covered by NHK News, one of Japan's leading national media outlets.	Remote <i>Dec 2025 - Feb 2026</i>
Research Intern, Prof. Hao Wang - Stevens Institute of Technology <i>Personalized federated learning, differential privacy, PyTorch/TensorFlow</i> Built personalized federated learning systems with differential privacy guarantees in PyTorch/TensorFlow to protect local data while accelerating on-device model personalization.	Hoboken, NJ <i>Jun 2024 - Feb 2026</i>
Software Engineering Intern, Geopogo <i>AI rendering platform, full-stack production engineering</i> - Shipped production AI architectural rendering platform as sole engineer, implementing React/TypeScript frontend, FastAPI backend, Firebase auth, and Vercel CI/CD. - Integrated image-reasoning and RunwayML text-to-video workflows with drag-and-drop uploads, real-time previews, and credit-based subscriptions.	Berkeley, CA <i>Sep 2025 - Dec 2025</i>

PROJECTS

Context8 - context8.org Built solo self-evolving Stack Overflow for AI agents where community votes and feedback turn one agent's solved bug into reusable knowledge for all agents; adjacent to later Context Hub and Evomap directions.	Oct 2025 - Jan 2026
AI Mario Level Generator - mario-hackmit.vercel.app <i>HackMIT 2025 - Modal Sponsor Prize</i> Built sketch-to-playable-level pipeline using LLaVA 1.5 for sketch interpretation, H100 GPU inference on Modal, FastAPI backend, React frontend, and OpenCV integration.	Sep 2025
Stud.ai <i>HackMIT 2024 - Smartest AI Agent Award</i> Built Chrome extension and AI agent that turns assignment rubrics into step-by-step timelines and auto-schedules work blocks on students' calendars.	Sep 2024

PUBLICATIONS

Agents' Last Exam (ALE) - arXiv:2606.05405v1 Co-first author. Benchmark of 1,000+ real-world agent tasks across 55 sub-fields and 13 industry clusters, led by 300+ industry experts and Prof. Dawn Song, UC Berkeley. arxiv.org/abs/2606.05405v1	Jun 2026
---	-----------------

AWARDS

Modal Sponsor Prize, HackMIT 2025; Smartest AI Agent Award, HackMIT 2024; Best Use of .Tech Domain, HackDavis 2024; Winner, Hack for Impact 2024.