

Yuanbo Pang

yuanbopang@berkeley.edu | 408-460-7526 | Berkeley, CA | linkedin.com/in/yuanbopang | github.com/PangYuanbo | yuanbopang.pro

UC Berkeley EECS student building AI agent evaluation, infrastructure, and robotics/vision systems; co-first author on the NeurIPS 2026 Agent's Last Exam benchmark submission.

SKILLS

Languages: Python, TypeScript/JavaScript, C++, Java, Rust | **Backend/Infra:** FastAPI, React, Node.js, Docker, Linux, Git, REST APIs, SQL, Firebase, Vercel, Modal, CI/CD | **AI Systems:** agent harnesses, benchmark infrastructure, Gemini API, LLaVA, PyTorch, TensorFlow | **Vision:** OpenCV, optical flow, 3D perception

EDUCATION

University of California, Berkeley <i>B.S. Electrical Engineering and Computer Sciences (EECS)</i> Coursework: CS 61B, CS 70, CS 170, EECS 16A, EECS 106A.	Berkeley, CA <i>Aug 2025 - Dec 2026</i>
Stanford University, Summer Session <i>CS 229: Machine Learning</i>	Stanford, CA <i>Summer 2025</i>
Community College Coursework <i>Computer Science coursework at De Anza, Foothill, Evergreen, and Diablo Valley; Phi Theta Kappa; Academic Director, Quantum Club; President, Data Club</i>	Cupertino, CA <i>Sep 2023 - Jun 2025</i>

EXPERIENCE

Full-Stack Engineer, Ludus <i>AI agent marketplace, full-stack systems, agent execution</i> - Built outcome-first marketplace where multiple coding/research agents execute the same task in parallel and users pay only for accepted results. - Built React/TypeScript frontend, Node/Hono backend, Neon Postgres task store, and Docker-based VM dispatch layer for agent execution.	Remote <i>Apr 2026 - Present</i>
Software Engineering Intern, Geopogo <i>AI rendering platform, full-stack production engineering</i> Shipped production AI rendering platform as sole engineer, implementing React/TypeScript frontend, FastAPI backend, Firebase auth, Vercel CI/CD, Gemini API, and RunwayML pipelines.	Berkeley, CA <i>Sep 2025 - Dec 2025</i>
ENGR Team Co-Lead, UC Berkeley BAIR - Agent's Last Exam (ALE) <i>Benchmark infrastructure, VM workflows, evaluation systems</i> - Built infrastructure for reproducible agent benchmark runs across VM-based GUI/CLI workflows, task setup, artifact capture, and deliverable-based scoring. - Converted expert-submitted tasks into executable benchmark instances spanning 55 sub-fields and 13 industry clusters. - Co-first author on ALE, a NeurIPS 2026 benchmark submission for evaluating frontier AI agents on long-horizon professional workflows.	Berkeley, CA <i>Jan 2026 - Present</i>
Robotics AI Engineering Intern, Starbot <i>Agent harness for restaurant service robots</i> - Built agent harness for restaurant service robots with tool APIs, task-state orchestration, and recovery paths for ordering, delivery, and customer-service workflows. - Designed tool interfaces, task state handling, and recovery paths to make real service workflows reliable enough for live demo settings. - System was demonstrated at CES 2026 and covered by NHK News, one of Japan's leading national media outlets.	Remote <i>Dec 2025 - Feb 2026</i>
Undergraduate Researcher, UC Berkeley - Prof. Avidoh Zakhor <i>UAV perception, optical flow, small-object detection</i> - Benchmarked UAV detection/tracking methods against 1,600 manually annotated DJI frames. - Evaluated optical-flow and background-compensation pipelines for tiny-object tracking and low-false-positive perception.	Berkeley, CA <i>Mar 2026 - Present</i>
Research Intern, Prof. Hao Wang - Stevens Institute of Technology <i>Personalized federated learning, differential privacy, PyTorch/TensorFlow</i> Built personalized federated learning systems with differential privacy guarantees in PyTorch/TensorFlow to protect local data while accelerating on-device model personalization.	Hoboken, NJ <i>Jun 2024 - Sep 2026</i>

PROJECTS

Context8 - context8.org <i>Agent knowledge systems, full-stack product</i> Built solo self-evolving Stack Overflow for AI agents with feedback loops that turn solved bugs into reusable knowledge for future agents.	<i>Oct 2025 - Jan 2026</i>
AI Mario Level Generator - mario-hackmit.vercel.app <i>HackMIT 2025 - Modal Sponsor Prize</i> Built GPU-backed sketch-to-game pipeline with LLaVA 1.5, Modal H100 inference, FastAPI, React, and OpenCV.	<i>Sep 2025</i>
Stud.ai <i>HackMIT 2024 - Smartest AI Agent Award</i> Built Chrome extension and AI scheduling agent for assignment planning workflows.	<i>Sep 2024</i>

PUBLICATIONS

Agent's Last Exam (ALE) - NeurIPS 2026 Evaluations and Datasets Track submission Co-first author. Benchmark of 1,000+ real-world agent tasks across 55 sub-fields and 13 industry clusters, led by 300+ industry experts and Prof. Dawn Song, UC Berkeley.	May 2026
---	-----------------

AWARDS

Modal Sponsor Prize, HackMIT 2025; Smartest AI Agent Award, HackMIT 2024; Best Use of .Tech Domain, HackDavis 2024; Winner, Hack for Impact 2024; First Prize, Jiangsu Province Chinese National Physics Olympiad 2022.