

Yuanbo Pang

yuanbopang@berkeley.edu | 408-460-7526 | San Jose, CA | linkedin.com/in/yuanbopang | github.com/PangYuanbo | yuanbopang.pro

EDUCATION

University of California, Berkeley <i>B.S. Electrical Engineering and Computer Sciences (EECS), Aug 2025 - Dec 2026</i> Coursework: CS 61B, CS 70, CS 170, CS 184, EECS 16A, EECS 106A. Stanford Summer School 2025: CS 229 Machine Learning.	Berkeley, CA
De Anza Community College <i>Computer Science transfer coursework, Sep 2023 - Jun 2025</i>	Cupertino, CA

SKILLS

Product/Full Stack: React, Next.js, TypeScript, FastAPI, Node.js/Hono, REST APIs, GraphQL, WebSockets
Systems: PostgreSQL/Neon, Redis, Firebase, AWS, Cloudflare, Railway, Vercel, CI/CD
AI Product: agent workflows, Claude Code, Codex, Cursor, Grok, Droid, LLaVA, OpenCV, model-backed interfaces, benchmark design

EXPERIENCE

Full-Stack Engineer, Ludus <i>AI agent marketplace, backend systems, full-stack product engineering</i> - Built outcome-first marketplace where multiple coding and research agents execute the same task in parallel and users pay for accepted results. - Implemented React/TypeScript frontend, Node/Hono backend, Neon Postgres task store, and Docker-based VM dispatch layer.	Apr 2026 - Present
Software Engineering Intern, Geopogo <i>Production AI rendering platform, sole engineer</i> Shipped React/TypeScript frontend, FastAPI backend, Firebase auth, Vercel CI/CD, image reasoning, and RunwayML video workflows.	Sep 2025 - Dec 2025
ENGR Team Co-Lead, UC Berkeley BAIR - Agent's Last Exam <i>Agent benchmark infrastructure and evaluation systems</i> Built reproducible benchmark workflows for long-horizon AI-agent tasks, including VM setup, artifact capture, and deliverable-based scoring.	Jan 2026 - Present

PROJECTS

Context8 - context8.org <i>Agent knowledge system</i> Built solo self-evolving Stack Overflow for AI agents where solved bugs become reusable knowledge through voting and feedback loops.	Oct 2025 - Jan 2026
AI Mario Level Generator <i>HackMIT 2025 Modal Sponsor Prize</i> Built sketch-to-playable-level pipeline with LLaVA 1.5, Modal H100 inference, FastAPI, React, and OpenCV.	Sep 2025