

Yuanbo Pang

yuanbopang@gmail.com | 408-460-7526 | San Jose, CA | linkedin.com/in/yuanbopang | github.com/PangYuanbo | yuanbopang.pro

EDUCATION

University of California, Berkeley <i>B.S. Electrical Engineering and Computer Sciences (EECS), expected Dec 2026</i> Coursework: Data Structures, Algorithms, Machine Learning, Discrete Math, Linear Algebra, Robotics.	Berkeley, CA
Stanford University, Summer Session <i>CS 229: Machine Learning</i>	Stanford, CA

SKILLS

Product/Systems: product engineering, AI workflow design, user research instincts, technical communication
Engineering: Python, TypeScript, React, FastAPI, SQL, REST APIs, Firebase
AI: agent evaluation, Gemini API, LLaVA, benchmark design

EXPERIENCE

Full-Stack Engineer, Ludus <i>AI agent marketplace, backend systems, full-stack product engineering</i> - Built outcome-first marketplace where multiple coding and research agents execute the same task in parallel and users pay for accepted results. - Implemented React/TypeScript frontend, Node/Hono backend, Neon Postgres task store, and Docker-based VM dispatch layer.	Apr 2026 - Present
Software Engineering Intern, Geopogo <i>Production AI rendering platform, sole engineer</i> Shipped React/TypeScript frontend, FastAPI backend, Firebase auth, Vercel CI/CD, Gemini image reasoning, and RunwayML video workflows.	Sep 2025 - Dec 2025
ENGR Team Co-Lead, UC Berkeley BAIR - Agent's Last Exam <i>Agent benchmark infrastructure and evaluation systems</i> Built reproducible benchmark workflows for long-horizon AI-agent tasks, including VM setup, artifact capture, and deliverable-based scoring.	Jan 2026 - Present

PROJECTS

Context8 - context8.org <i>Agent knowledge system</i> Built solo self-evolving Stack Overflow for AI agents where solved bugs become reusable knowledge through voting and feedback loops.	Oct 2025 - Jan 2026
AI Mario Level Generator <i>HackMIT 2025 Modal Sponsor Prize</i> Built sketch-to-playable-level pipeline with LLaVA 1.5, Modal H100 inference, FastAPI, React, and OpenCV.	Sep 2025