

Yuanbo Pang

yuanbopang@berkeley.edu | 408-460-7526 | Berkeley, CA | linkedin.com/in/yuanbopang | github.com/PangYuanbo | yuanbopang.pro

UC Berkeley EECS student building AI agent evaluation, infrastructure, and robotics/vision systems; co-first author on the NeurIPS 2026 Agent's Last Exam benchmark submission.

SKILLS

Languages: Python, TypeScript/JavaScript, C++, Java, Rust, SQL | AI/Agents: agent harnesses, benchmark evaluation, LLaVA, Gemini API, PyTorch, TensorFlow | Robotics/Vision: 3D perception, SLAM, UAV detection, optical flow, OpenCV | Infra: FastAPI, React, Node.js, Docker, Linux, Git, REST APIs, CI/CD, Firebase, Vercel, Modal

EDUCATION

University of California, Berkeley
B.S. Electrical Engineering and Computer Sciences (EECS)
• Coursework: CS 61B, CS 70, CS 170, EECS 16A, EECS 106A.

Stanford University, Summer Session
CS 229: Machine Learning

Community College Coursework
Computer Science coursework at De Anza, Foothill, Evergreen, and Diablo Valley; Phi Theta Kappa; Academic Director, Quantum Club; President, Data Club

Berkeley, CA
Aug 2025 - Dec 2026

Stanford, CA
Summer 2025

Cupertino, CA
Sep 2023 - Jun 2025

EXPERIENCE

Full-Stack Engineer, Ludus
AI agent marketplace, full-stack systems, agent execution
• Built outcome-first marketplace where multiple coding/research agents execute the same task in parallel and users pay only for accepted results.
• Built React/TypeScript frontend, Node/Hono backend, Neon Postgres task store, and Docker-based VM dispatch layer for agent execution.

ENGR Team Co-Lead, UC Berkeley BAIR - Agent's Last Exam (ALE)
AI agent evaluation, benchmark infrastructure, agent harnesses
• Co-first author on ALE, a NeurIPS 2026 benchmark submission led by Prof. Dawn Song for evaluating frontier AI agents on long-horizon professional workflows.
• Primary builder for turning expert-submitted real-world tasks into executable benchmark instances across 55 sub-fields and 13 industry clusters.
• Built agent evaluation infrastructure for VM-based GUI/CLI workflows, reproducible harness runs, and deliverable-based scoring.

Software Engineering Intern, Geopogo
AI rendering platform, full-stack production engineering
• Shipped production AI architectural rendering platform as sole engineer, implementing React/TypeScript frontend, FastAPI backend, Firebase auth, and Vercel CI/CD.
• Integrated Gemini image reasoning and RunwayML text-to-video workflows with drag-and-drop uploads, real-time previews, and credit-based subscriptions.

Robotics AI Engineering Intern, Starbot
Agent harness for restaurant service robots
• Built agent harness for restaurant service robots with tool APIs, task-state orchestration, and recovery paths for ordering, delivery, and customer-service workflows.
• System was demonstrated at CES 2026 and covered by NHK News, one of Japan's leading national media outlets.

Undergraduate Researcher, UC Berkeley - Prof. Avidesh Zakhor
UAV small-object detection, motion priors, airborne target tracking
• Research small UAV detection and tracking in aerial video, with focus on tiny-object recognition, camera-motion compensation, and low-false-positive perception for onboard drone systems.
• Benchmarked UAV detection methods including TransVisDrone, AOT Winner v022, ESOD, and Li-TETC/NPS against 1,600 manually annotated DJI frames.
• Evaluated optical-flow and background-compensation pipelines to reduce false positives in tiny airborne object tracking.

Research Intern, Prof. Hao Wang - Stevens Institute of Technology
Personalized federated learning, differential privacy, PyTorch/TensorFlow
• Built personalized federated learning systems with differential privacy guarantees in PyTorch/TensorFlow to protect local data while accelerating on-device model personalization.

Remote
Apr 2026 - Present

Berkeley, CA
Jan 2026 - Present

Berkeley, CA
Sep 2025 - Dec 2025

Remote
Dec 2025 - Feb 2026

Berkeley, CA
Mar 2026 - Present

Hoboken, NJ
Jun 2024 - Feb 2026

PROJECTS

Context8 - context8.org
• Built solo self-evolving Stack Overflow for AI agents where community votes and feedback turn one agent's solved bug into reusable knowledge for all agents; adjacent to later Context Hub and Evomap directions.

AI Mario Level Generator - mario-hackmit.vercel.app
HackMIT 2025 - Modal Sponsor Prize
• Built sketch-to-playable-level pipeline using LLaVa 1.5 for sketch interpretation, H100 GPU inference on Modal, FastAPI backend, React frontend, and OpenCV integration.

Stud.ai
HackMIT 2024 - Smartest AI Agent Award
• Built Chrome extension and AI agent that turns assignment rubrics into step-by-step timelines and auto-schedules work blocks on students' calendars.

Oct 2025 - Jan 2026

Sep 2025

Sep 2024

PUBLICATIONS

Agent's Last Exam (ALE) - NeurIPS 2026 Evaluations and Datasets Track submission
• Co-first author. Benchmark of 1,000+ real-world agent tasks across 55 sub-fields and 13 industry clusters, led by 300+ industry experts and Prof. Dawn Song, UC Berkeley.

May 2026

AWARDS

Modal Sponsor Prize, HackMIT 2025; Smartest AI Agent Award, HackMIT 2024; Best Use of .Tech Domain, HackDavis 2024; Winner, Hack for Impact 2024; First Prize, Jiangsu Province Chinese National Physics Olympiad 2022.